2026 Global IT Infrastructure and Geopolitical Risk Assessment Report: Verification of Microsoft's Platform Strategy and Macroeconomic Ripple Effects

Introduction and Executive Summary

As of 2026, the global technology ecosystem and macroeconomy are in the midst of a polycrisis driven by unprecedented capital expenditures in artificial intelligence (AI) infrastructure, the resulting depletion of physical resources, and a catastrophic geopolitical shock in the Middle East. This report rigorously verifies the claims made in the provided document, "Will you still continue to use Windows and Office?", against the latest market data, academic research on computational complexity, macroeconomic indicators, and geopolitical dynamics to extract and analyze the underlying ripple effects.

The analysis confirms that Microsoft's strategy of deeply integrating "Copilot" into its operating system (OS), progressively raising hardware requirements, and exercising monopolistic consumption of the energy and semiconductor markets as a hyperscaler has mutated into a structure that forces unauditable compliance risks and astronomical costs upon enterprises. Microsoft's platform, once regarded as "neutral and stable infrastructure supporting business operations," currently functions as a high-risk agent that strips users of their autonomy and fundamentally distorts global supply chains. This report details the outcome of the enterprise AI market battle, the mathematical limitations of the Transformer architecture, the e-waste crisis, the regression to fossil fuels, and the realities of resource triage brought about by the Iran War that erupted in February 2026, offering strategic recommendations for Chief Information Officers (CIOs) and Chief Information Security Officers (CISOs) of enterprises and public institutions.

Seismic Shifts in the Enterprise AI Market and Copilot's Structural Defeat

The subject document's primary claim that "the outcome in the enterprise market is already decided" is fully corroborated by Q1 2026 market research data and the financial metrics of the respective companies. Despite possessing an historically unprecedented distribution channel advantage by bundling Copilot as a standard feature in its OS and Office applications, Microsoft is facing severe rejection and stagnant usage rates in the active market. This phenomenon represents a paradigm shift that overturns the software industry's long-standing dogma that

"control of distribution channels guarantees market hegemony."

Disconnect in Adoption Rates and the Collapse of Net Promoter Scores

The "AI Choice 2026" large-scale survey conducted by Recon Analytics in January 2026, targeting over 150,000 paid AI subscribers in the U.S., highlights a severe contraction in Microsoft Copilot's market position.[11] According to the data, Copilot's market share among U.S. paid AI subscribers plummeted from 18.8% in July 2025 to 11.5% in January 2026.[1] This represents a massive 39% contraction in market position over just seven months, and it currently struggles to stabilize in the 10% to 12% range.[11]

The root cause of this market share loss lies in the massive disconnect between provisioned seats and actual active usage by employees. Of the employees granted access to Copilot in the workplace, only 35.8% actually convert to using it continuously as an active tool.[11] This figure indicates that roughly two-thirds of employees at companies that purchased the license actively exclude Copilot from their daily workflows despite having access to it.

The data becomes even more revealing in environments where employees are given freedom of choice. When an enterprise provides Copilot as the sole platform, 68% of employees adopt it as their primary tool; however, when alternative platforms like ChatGPT or Google Gemini are provided simultaneously, Copilot's adoption rate drops precipitously.[11] In environments where both ChatGPT and Copilot are available, Copilot's adoption falls to 18%, and in environments where all three major platforms (Copilot, ChatGPT, Gemini) are available, the number of users choosing Copilot collapses to a mere 8%.[11] This objectively proves that the unparalleled accessibility of being embedded by default in the OS and Office apps does not translate into trust in product quality or actual usage.

An even more fatal metric is the Net Promoter Score (NPS), which indicates product recommendation and satisfaction. The NPS for Copilot's accuracy has remained persistently negative since recording -3.5 in July 2025, plunging to -24.1 in September 2025, and still showing a severe level of distrust at -19.8 as of January 2026.[11] Copilot's satisfaction score regarding accuracy is 23 points below market leader Google Gemini and 9 points below ChatGPT.[11] Among users who completely stopped using Copilot, 44.2% cited "distrust of answers" as their primary reason for leaving, exhibiting a prominent rejection response compared to competing platforms.[11]

Anthropic (Claude)'s Explosive Growth and Enterprise Market Monopoly

In stark contrast to Microsoft Copilot's stagnation, Anthropic (Claude)'s establishment of hegemony in the enterprise AI market became a definitive reality through a series of financial disclosures in the spring of 2026. Anthropic has charted the fastest growth trajectory in history for a B2B software company, effectively monopolizing enterprise market budgets.

Financial & Adoption Metrics	Anthropic (Claude) Reality	Macro Trends & Implications
ARR (Annual Recurring Revenue)	Surpassed $30 billion (April 2026) [23]	Overtook OpenAI's estimated $24B-$25B to become the sole industry leader.[42]
Revenue Growth Rate	Approx. 30x in the past 15 months (from $1B in Jan 2025) [52]	Surpassed the growth speeds of historic SaaS companies like Slack, Zoom, and Snowflake.[2]
Corporate Valuation (Private)	$380 billion (Feb 2026 Series G) [46]	Completed $30B funding in Series G from GIC, Coatue, etc..[67]
Secondary Market Valuation	Approx. $1 trillion (Forge Global, etc.) [45]	Trading at 2.6x primary price, reversing OpenAI (approx. $880B).[4]
Large Enterprise Adoption	Customers include 70% of Fortune 100, 8 of Fortune 10 [58]	Exceeded 1,000 large enterprise customers spending over $1 million annually.[53]
Developer Tool Revenue	Claude Code alone reached $2.5B ARR [48]	Proof of overwhelming reasoning capability in code generation and autonomous agent functions.[4]

A defining characteristic of Anthropic's revenue structure is that the overwhelming majority (70-75%) consists of enterprise API usage and corporate subscriptions.[8] While OpenAI leans heavily toward consumers with hundreds of millions of free users, Anthropic is intensively raising capital from Fortune 100 companies that demand strict data governance and hallucination resistance.[538]

The secondary and tertiary insights implied by this contrast are vital. The market is not rejecting the concept of "AI" itself. Enterprises are ungrudgingly allocating massive budgets in the millions of dollars for "trustworthy architectures" capable of handling complex logical reasoning and secure coding tasks.[53] On the other hand, even if Microsoft's Copilot is bundled for free or with low friction within the existing Microsoft 365 ecosystem, enterprises are passing on its adoption,

viewing it as a system incapable of guaranteeing practical complexity. The fact that in a Gartner survey, only 24% of companies that piloted Copilot were able to roll it out to more than 20% of their employees [9] further corroborates the structural limitations of the product.

While Microsoft is pouring an astronomical $100 billion to $120 billion in capital expenditure (Capex) into infrastructure [10], the fact that Copilot—running on that massive hardware—is being pushed out by market choice raises fundamental doubts about the viability of recouping those investments.

Computational Complexity of the Transformer Architecture and the Mathematical Inevitability of Hallucinations

The fundamental reason enterprises harbor deep distrust toward the output of Copilot and general Large Language Models (LLMs) (as seen in the NPS collapse) is not merely due to biased training data or temporary software bugs. It stems from the mathematical and physical computational limitations inherent in the foundational Transformer architecture itself. This essential flaw has been rigorously proven from the perspective of computational complexity theory by researchers at Stanford University.

Sikka & Sikka (2025) Theorem and the "Silent Breakdown"

In July 2025, a paper titled "Hallucination Stations: On Some Basic Limitations of Transformer-Based Language Models" (arXiv:2507.07505) published by Varin Sikka and Vishal Sikka (former SAP CTO) at Stanford University proved that hallucinations (the generation of plausible but false information) caused by LLMs are an unavoidable mathematical phenomenon caused by a computational complexity mismatch.[11][12][13]

The core theorem of this paper is based on the fact that the inference process (Self-Attention mechanism) used by a standard Transformer-based LLM to generate text operates under a strict "fixed budget" constraint, always executing in $\mathcal{O}(N^2 \cdot d)$ computation time relative to the number of input tokens N and the vector dimension d.[11][12][13] Even if the model's context window (N) expands to millions of tokens, the total volume of floating-point operations allocated for generating each token is geometrically fixed by the model's architecture.[11]

The problem is that many of the practical tasks enterprises demand of LLMs possess an algorithmic complexity that asymptotically far exceeds this $\mathcal{O}(N^2 \cdot d)$ compute budget.[11][14]

- Large-scale matrix multiplications or relational algebraic joins across multiple databases (such as lengthy Excel sheets) require at least $O(n^3)$ compute time.[11][14]
- Tasks strictly verifying the solutions to NP-hard scheduling or optimization problems,

typified by the Traveling Salesperson Problem (TSP), require computational steps on a factorial order of $(n-1)!/2$.[1111]

According to Sikka et al.'s theorem, the moment the inherent computational complexity of a task requested by a prompt exceeds the model's inference limit of $\mathcal{O}(N^2 \cdot d)$, it becomes physically and mathematically impossible for the LLM to execute the correct computational steps for the task.[1112] However, the most fatal flaw is that LLMs lack an explicit error declaration function (fail-safe) to state, "processing interrupted due to exceeding computational capacity." An LLM that hits a computational wall merely falls back to a probabilistic mode, blindly generating "the next token deemed most contextually probable".[11]

The result is a "confident incorrect answer" that is grammatically perfect but logically entirely broken—namely, a hallucination. This is not a bug; it is an unavoidable consequence driven by the architecture's specifications.[1112]

The Impossibility of Verification in Agentic AI and the Compliance Crisis of OS Integration

The paper further notes that autonomously operating "LLM agents" cannot escape the curse of this computational limit.[11] The task for an AI agent (Agent B) to verify whether the solution generated by another AI agent (Agent A) is correct often entails computational complexity equal to or higher than generating the solution itself.[1114] Therefore, even if multiple agents are combined, as long as the underlying inference model shares the same $\mathcal{O}(N^2 \cdot d)$ limitation, it is theoretically impossible to guarantee system accuracy.[1114]

In light of this mathematical limitation, it becomes apparent how dangerous Microsoft's strategy is to integrate a system harboring this essential flaw (Copilot) into the deepest layers of the Windows OS and Microsoft 365 by default. Whether editing a long contract in Word, running complex financial aggregations in Excel, or summarizing confidential emails in Outlook, the moment task complexity crosses the threshold, the system triggers a "silent breakdown."

From the user's perspective, there is no external way to distinguish between a result where the system correctly completed the computation and a hallucination generated probabilistically after exceeding compute limits. Through this unauditable process, fabricated financial transactions, incorrect medical prescription summaries, or inaccurate administrative documents created by AI accumulate within the system without verbatim human verification, materializing the risk of "Shadow AI." Forcing this architecture—which obscures the locus of legal liability—must be deemed an artificial injection of an unprecedented cybersecurity vulnerability into every critical infrastructure requiring strict compliance, such as healthcare, finance, and government.

OS Tyranny and the Forced Hardware Lifecycle

Under the pretext of AI integration, Microsoft has arbitrarily manipulated OS requirements,

forcing a coercive hardware replacement cycle onto end users. This approach imposes a massive financial burden on the global economy while simultaneously causing irreversible environmental damage.

The High Cost Structure of Windows 10 ESU and the Increase in TCO

With the end of support for Windows 10 scheduled for October 14, 2025, purchasing the Extended Security Updates (ESU) becomes unavoidable to maintain OS security.[1516] The pricing structure for enterprise ESUs designed by Microsoft employs a punitive, progressive scale intended to force users onto the new OS.

Coverage Period	Annual Cost Per Device	Cumulative Total Cost of Ownership (TCO)
Year 1 (Nov 2025 - Oct 2026)	$61 [1517]	$61
Year 2 (Nov 2026 - Oct 2027)	$122 (Double previous year) [1517]	$183
Year 3 (Nov 2027 - Oct 2028)	$244 (Double previous year) [1517]	$427 [1518]

This program is cumulative, meaning that even if an enterprise attempts to join from Year Two, they must pay the costs retroactively for Year One.[1517] Furthermore, ESUs provide patches only for severe security vulnerabilities and explicitly exclude the addition of new features, performance improvements, or general technical support.[1517] Enterprises are forced to choose between paying a massive "ransom" of $427 per device over three years simply to "maintain" their existing infrastructure, or purchasing entirely new Copilot+ PCs amidst a soaring memory market.

Increased Requirements for Windows 11 24H2 and the Massive Generation of E-Waste

The coercion of hardware updates is also relentlessly executed in the transition to Windows 11. Upon the release of Windows 11 in 2021, Microsoft mandated TPM 2.0 and Intel 8th Gen/Ryzen 2000-series or later processors, cutting off PCs from 2017 and earlier. Furthermore, starting with the "24H2" version released in 2024, CPU instruction set requirements were covertly tightened, making the "SSE4.2" and "POPCNT" instructions mandatory.[1920]

This physically blocked the methods IT administrators previously used to extend the life of older hardware, such as bypassing the TPM check using registry hacks or installation media creation

tools like Rufus.[19] If the CPU lacks these instruction sets (e.g., the Intel Core 2 Duo generation), the OS kernel itself refuses to boot.[19][20]

According to rigorous analysis by research firm Canalys, this series of raised hardware requirements combined with the end of Windows 10 support seals the fate of approximately 240 million PCs—equivalent to about 20% of Windows devices globally—to be discarded as electronic waste (e-waste) "even though they are fully operational".[21][22][23]

The total weight of these discarded PCs reaches an astronomical 480 million kilograms, equivalent to 320,000 cars.[21][24][22] The vast majority of these discarded PCs head straight for landfills, where hazardous materials like lead, mercury, cadmium, and flame retardants will cause severe long-term contamination of soil and groundwater. Many of these 240 million devices are perfectly sound hardware and could easily be used for several more years in educational settings or basic administrative work if a lightweight open-source OS like Linux were installed. Microsoft's strategy to forcibly obsolete functionally sound hardware for the singular corporate goal of AI integration is a grave betrayal of global environmental sustainability.

Deprivation of User Sovereignty and Dark Patterns

Microsoft's push for Copilot integration goes beyond forcing hardware replacements; it violently alters the operating system's interface itself, stripping users of their autonomy. This transition to an "un-opt-outable AI ecosystem" has drawn strong condemnation from industry experts.

Criticism from Mozilla and the Reality of Forced Automatic Installation

In April 2026, Mozilla VP Linda Griffin publicly criticized Microsoft's heavy-handed software deployment methods in a statement titled "Old habits die hard: Microsoft tries to limit our options, this time with AI".[25][26][26]

The primary issue Griffin highlighted was the fact that the M365 Copilot app was silently auto-installed in the background on Windows devices running Microsoft 365 desktop apps without any prior prompts or explicit user consent (opt-in).[25][26] This tactic of embedding an AI tool deep within the OS without the user's knowledge, where it begins collecting and learning data behind the scenes, was condemned as "a typical dark pattern prioritizing the company's profits by stripping users of their choices".[25][26]

Moreover, the design stance of forcing the placement of a dedicated key to launch Copilot on physical keyboards, while making it technically difficult to remap the function of that key for other purposes, is an infringement on users' rights to control their own devices (user sovereignty).[26]

It is particularly noteworthy that only users in the European Economic Area (EEA) have been excluded from this automatic Copilot installation due to strict antitrust and data protection regulations like the Digital Markets Act (DMA). This proves that Microsoft's management acknowledges this aggressive installation method is legally problematic and contrary to user interests. Enterprises and users in Japan, the U.S., Asia, Latin America, and Africa, who are unprotected by such regulations, are placed under arbitrary platform surveillance with complete

disregard for corporate ethics.

Macroeconomic Crowding-Out Effects by Hyperscalers

To prop up unused AI and Copilot with its stagnating adoption rate, Microsoft is pouring astronomical capital expenditures (Capex)—between $100 billion and $146 billion annually, rivaling the national budgets of many mid-sized nations—into data centers and semiconductor infrastructure.[10] This abnormal scale of capital injection has transcended mere "investment" in a free market, triggering a macroeconomic "crowding-out effect" that drains resources from other industries.

Abnormal Surge in Semiconductor Memory Prices (DRAM and NAND Shortages)

The hyperscalers' frantic investments are causing the most direct destruction in the semiconductor memory market. AI servers for data centers necessitate High Bandwidth Memory (HBM) and ultra-high-capacity enterprise SSDs to support massive parallel processing. To meet long-term supply contracts from the likes of Microsoft and OpenAI, giant semiconductor manufacturers like Samsung and SK Hynix, which command about 78% of the global share, dramatically shifted their production lines away from general-purpose memory toward high-value-added memory for AI and servers.

As a result of this extreme reallocation of supply capacity, the market supply of general-purpose DRAM and NAND flash memory used in smartphones, standard PCs, and consumer electronics has completely dried up, triggering a collapsing hyperinflation of prices.

Market data from Counterpoint Research captures this disastrous situation in detail.[272829]

Metric / Period	Counterpoint Research Price Trend Data
DRAM Price Surge (Q3 2025)	Recorded a historic 172% year-over-year increase.[27]
DRAM Price Surge (Q1 2026)	Unprecedented 80-90% quarter-over-quarter surge.[28]
64GB RDIMM Price Trend	Jumped from $255 in Q3 2025 to $450 in Q4 2025, reaching $700 by March 2026.[29] This is nearly double the 2018 peak price.[29]
Profit Margin Inversion	In Q4 2025, the operating margin of

	general-purpose DRAM surpassed HBM for the first time in history, reaching the 60% range.[28] A complete seller's market has emerged.

This sharp increase in Bill of Materials (BOM) costs is directly hitting device manufacturers. PC and smartphone makers such as HP, Lenovo, Xiaomi, and Dell are unable to absorb the soaring costs through corporate efforts alone, forcing them to raise retail prices by 15% to 25%. Counterpoint Research analysts warn that "For device manufacturers, this is a double whammy – rising component costs and weakened consumer purchasing power will likely slow the demand".[28]

General consumers not using AI, small-to-medium enterprises with limited infrastructure budgets, and IT procurement for educational institutions are all effectively being forced to pay an "invisible tax to subsidize Microsoft's AI infrastructure buildup" in the form of hardware price markups. This is the macroeconomic reality of 2026.

Data Center Construction Delays and the Scramble for Physical Materials

The crowding-out effect has also formed severe bottlenecks in electrical infrastructure and physical materials. Interconnection queues to the U.S. power grid have swelled past 2,100 GW, and delivery times for the massive transformers required to supply power to data centers have stretched up to 120 weeks (over two years), with prices soaring four to six times compared to 2022.

Consequently, only about a third of the data center capacity scheduled to go online in the U.S. in 2026 is actually under construction, with 30% to 50% of the remainder delayed to 2028 or later due to electrical equipment delivery backlogs. A further severe ripple effect is that as data center supply chains prioritize Grain-Oriented Electrical Steel (GOES) for transformer cores and copper essential for AI server cooling and wiring, local U.S. utility companies are forced to delay replacement programs for aging distribution transformers, structurally elevating blackout risks and grid vulnerability in general society.

Distortions in Energy Markets and the Hollowed-Out Climate Change Measures

To secure power for expanding data centers, Microsoft has abandoned the illusion of clean growth via renewable energy and pivoted sharply toward massive direct investments in fossil fuel infrastructure. The ambitions of a single tech giant are setting back global climate change mitigation progress by decades.

The $7 Billion Deal with Chevron and Engine No. 1

In April 2026, Microsoft entered negotiations with oil and gas major Chevron and investment

firm Engine No. 1 to forge a $7 billion exclusive partnership for the joint development of a massive natural gas-fired power plant in the Permian Basin of West Texas.[30313210]

This project is slated to generate 2,500 megawatts (2.5 GW) of electricity initially, with plans to expand to up to 5,000 megawatts (5.0 GW) in the future.[303133] Powering 5.0 GW is equivalent to the annual consumption of approximately 4 million standard American households.[31]

A primary feature of this power plant is that, to avoid straining the already fragile regional power grid (and avoiding grid connection queue delays), it employs a "behind-the-meter" method, utilizing abundant natural gas from the Permian Basin to supply electricity directly to an adjacent 7,000-acre Microsoft data center campus.[3430] For Chevron, this implies not merely wholesale fuel sales, but a powerful new business model of converting its gas into "high-value-added AI electricity" and selling it directly.[31]

Collapse of the "Carbon Negative Pledge" and Industry-Wide Emissions Increases

However, this return to fossil fuels solidifies the complete hollowing out of Microsoft's ambitious environmental pledge, declared loudly in 2020, to "achieve carbon negative by 2030."

According to the "FY2025 Annual Climate Report" released by the New York City Employees' Retirement System (NYCERS) Comptroller's Office, Microsoft's reported greenhouse gas (CO_2) emissions have already surged by 23.4% compared to 2020, driven primarily by explosive AI and cloud infrastructure demand.[3536] The report also highlights that Google recorded a 51% emissions increase compared to 2019 [36], indicating that the environmental burden of AI infrastructure is a structural issue across the entire tech industry.

Once the massive $7 billion natural gas plant in Texas comes online after 2027, the trajectory of these emissions increases is certain to become even steeper. Cheap natural gas resources, which should be utilized for household heating, industrial boilers, and general grid stabilization, are being preferentially diverted and burned at AI data centers simply to support Copilot, which suffers from low active usage.

National Sovereignty and the Geopolitics of Sovereign AI Infrastructure

Microsoft's investments extend far beyond the bounds of mere corporate activity, penetrating deeply into national security and data sovereignty domains of various countries. A symbolic case is its massive AI infrastructure investment in Japan.

In April 2026, Microsoft announced a $10 billion (approx. 1.6 trillion yen) total investment over the four years from 2026 to 2029 to bolster AI infrastructure, cybersecurity, and talent development in Japan.[373839] The core of this investment is the construction of a "Sovereign AI" environment that enables advanced AI processing while keeping data within Japan's

borders.[3738]

Microsoft established strategic alliances with Japan's major domestic telecommunications and cloud providers, SoftBank and Sakura Internet, to build a framework offering GPU-based AI computing platforms under the Microsoft Azure environment through their infrastructure.[373840] This operates fully in tandem with the multi-trillion-yen AI infrastructure plan pushed by Japan's Ministry of Economy, Trade and Industry (METI), acting as a scheme leveraging massive government subsidies.[37]

Furthermore, Microsoft has deepened collaborations with core national security agencies like Japan's National Police Agency (NPA) and the National center of Incident readiness and Strategy for Cybersecurity (NISC), advancing efforts to adapt Azure's security baseline to the Japanese threat landscape.[38] In addition, setting a goal to train 1 million AI professionals by 2030, it has incorporated Japan's legacy IT giants—Fujitsu, Hitachi, NEC, and NTT Data—into its ecosystem.[413842]

The Japanese government desired a "domestic-closed" infrastructure to avoid the risk of sensitive data crossing borders to be processed on overseas servers.[3738] By aptly responding to this political demand and injecting immense capital, Microsoft successfully achieved a "Sovereign Lock-in" of Japan's public and private IT infrastructure into its own architecture.[37]

However, for Japan, a country with an energy self-sufficiency rate below 10%, building a network of data centers that consume enormous power while functionally outsourcing the brain of the operation to an American private enterprise can hardly be called the establishment of true "sovereignty." Rather, it should be evaluated as a process of "dependent vassalization," rendering the country extremely vulnerable to energy supply disruptions and macroeconomic shocks.

Geopolitical Black Swan: The 2026 Iran War and the Collapse of Global Supply Chains

While hyperscalers like Microsoft push forward with $100 billion investment plans predicated on cheap energy and stable supply chains, a geopolitical "Black Swan" emerged in the Middle East at the end of February 2026, shattering all those assumptions.

Operation Epic Fury and the "Dual Blockade" of the Strait of Hormuz

On February 28, 2026, a joint force of the U.S. and Israel executed a massive surprise airstrike campaign, codenamed "Operation Epic Fury," targeting Iran's nuclear facilities, military infrastructure, and government centers.[434445] In the first 12 hours alone, nearly 900 precision-guided munitions and Tomahawk cruise missiles were fired [444645], resulting in the assassination of Iran's Supreme Leader Ayatollah Ali Khamenei, Secretary of the Supreme National Security Council Ali Larijani, and numerous other top government officials.[4347]

In retaliation for this devastating attack, Iran's Islamic Revolutionary Guard Corps (IRGC)

immediately launched counterstrikes against Israel and U.S. bases in the Middle East using missiles and drones, and took the drastic measure of completely cutting off navigation through the Strait of Hormuz, the world's major energy artery.[4348] The Strait of Hormuz is a strategic maritime chokepoint through which approximately 25% of the world's seaborne crude oil trade and 20% of its liquefied natural gas (LNG) trade pass; this blockade triggered the largest supply shock in the history of the world oil market, surpassing the oil crises of the 1970s.[4848]

Making the situation even more dire were the countermeasures taken by the U.S. Department of Defense. Under President Trump's orders, the U.S. Navy initiated a massive maritime blockade against major Iranian oil export ports (such as Kharg Island) starting around mid-April, forcefully turning back all ships calling at or departing from Iran.[434950] This resulted in a worst-case stalemate—a "Dual Blockade"—where Iran's physical closure of the strait and the U.S. Navy's economic blockade of Iran were occurring simultaneously.[4348]

Outbreak of the Naphtha Crisis and Ultimate Resource Triage

As Brent crude oil prices spiked to historic highs of $126 a barrel [48], and gasoline and electricity costs surged globally, a more profound and irreversible destruction silently progressed deep within the global supply chain: the "Naphtha Crisis."

Naphtha is a basic chemical feedstock derived from the crude oil distillation process, cracked to produce ethylene and propylene. These are irreplaceable baseline materials for virtually all products in modern society, including polyvinyl chloride (PVC), polyethylene (PE), polyurethane, synthetic fibers, and chemicals used in semiconductor manufacturing.

The dual blockade of the Strait of Hormuz dealt a catastrophic blow to the petrochemical industries in Asia, with Japan depending on Middle Eastern crude imports for over 95% and South Korea for over 70% of their supply.[48] Immediately following the outbreak of the war, naphtha spot prices in Asian markets nearly doubled from about $600 to $1,190 per ton, forcing major Japanese chemical manufacturers like Idemitsu Kosan and Tosoh to scale back production and implement massive price hikes.

It is here that Microsoft's frantic AI infrastructure investments and the geopolitical resource exhaustion collide in the cruellest fashion.

A scramble for plastics between medical infrastructure and AI infrastructure.

The construction and operation of data centers require enormous volumes of petrochemical products. High-quality PVC is used in the insulation of tens of thousands of kilometers of server cables; PE is essential for the specialized piping of modern immersion cooling systems and fiber optic cladding; and massive quantities of polyurethane and specialty resins are consumed in server chassis components.

Simultaneously, the modern medical system is completely dependent on these exact same basic chemicals. Sterile, disposable medical devices—such as dialyzers used in artificial dialysis, blood tubing, IV bags, syringes, and catheters—are entirely manufactured from PVC

and PE.

The extreme scarcity and skyrocketing price of naphtha triggered an ultimate "resource triage" in the market. Under normal circumstances, the price mechanism ensures that materials critical for life-saving medical infrastructure are prioritized. However, in the abnormal macroeconomic environment of spring 2026, hyperscalers like Microsoft, wielding bottomless capital expenditure budgets (Capex) exceeding $100 billion, secured raw materials (PVC, PE, etc.) via forward purchases at the highest upstream points of the supply chain—disregarding the inflated costs—simply to prevent delays in data center construction.

Consequently, medical device manufacturers and small-to-medium packaging companies, vastly inferior in financial clout, could not pass the multiplied material costs onto their product pricing and were physically locked out of the procurement market. The stark reality is that resources vital for sustaining the lives of dialysis patients are being ruthlessly siphoned off to manufacture cables for AI data centers that see an active usage rate of just 35.8%. This is not an anti-technology metaphor; it is the cold physical reality produced by supply chains operating under the Hormuz blockade.

Strategic Conclusions and Recommendations for Organizational Defense

Microsoft, led by Satya Nadella, placed a massive gamble on integrating a single AI product, Copilot, into the deep layers of social infrastructure. However, facing the "walls of reality"—market rejection, the exposure of the architecture's mathematical limits, and geopolitical depletion of physical resources—this strategy has entered a self-destructive negative loop.

The pressure to recoup massive investments, which cannot be halted, has culminated in forcefully raising hardware requirements for users (generating massive e-waste), embedding unconsented surveillance processes at the OS level, and relentlessly exploiting energy and fossil fuels. Microsoft's platform has relinquished its role as a "neutral infrastructure" for enhancing organizational productivity, mutating into a "high-risk agent" that endlessly drains client budgets and threatens compliance liabilities merely to justify its own stock price and infrastructure expenditures.

Strategic Recommendations for IT Leaders (CIOs / CISOs)

Chief Information Officers (CIOs) and Chief Information Security Officers (CISOs) of enterprises, municipalities, medical institutions, and financial organizations must face this macroeconomic and technological reality and immediately execute the following strategic pivots:

1. **Break Free from Vendor Lock-in and Restore Sovereignty:**
   Organizations must reject punitive payments for Windows 10 ESUs and massive capital expenditures on unnecessary Copilot+ PCs to reclaim their autonomy. Specifically, for departments handling sensitive data (medical records, financial audits, public services), IT

leaders should avoid the Windows 11 environment where unauditable AI processing pathways are open by default, and aggressively promote a phased transition to open-source OSes like Linux for desktop environments.

2. **Select Verifiable AI Architectures:** Rather than rejecting AI adoption altogether, organizations should adopt an approach that integrates the APIs of independent vendors—like Anthropic (Claude)—who have secured market trust (NPS) in enterprise-level strict data governance and complex reasoning tasks.[18] This integration should be strictly limited to a verified "opt-in" basis only for necessary processes.
3. **Re-evaluate Supply Chain Risks:**
   As the geopolitical crises (the Iran War and Naphtha Crisis) have proven, the monopolization of resources by hyperscalers drives hardware procurement costs (DRAM, etc.) to unpredictable levels. Over-reliance on a single cloud provider (e.g., Azure) or a specific hardware ecosystem constitutes a fatal vulnerability to an organization's Business Continuity Plan (BCP).

The most reliable and rational defense in technology strategy is to physically and logically detach organizational assets from the massive silos of runaway platformers embroiled in overlapping global crises. The professional answer to the question "Will you still continue to use Windows and Office?" is now a definitive "No." True security and productivity can only be guaranteed by transitioning to independent, self-sovereign architectures.

引用文献

1. AI Choice 2026: Why Licenses Don't Equal Adoption - Recon Analytics, 4月 26, 2026にアクセス、https://www.reconanalytics.com/ai-choice-2026-why-licenses-dont-equal-adoption/
2. Anthropic Just Passed OpenAI in Revenue. While Spending 4x Less to Train Their Models, 4月 26, 2026にアクセス、https://www.saastr.com/anthropic-just-passed-openai-in-revenue-while-spending-4x-less-to-train-their-models/
3. Anthropic Just Passed OpenAI in Revenue. Here Is Why It Matters. | by David C. - Medium, 4月 26, 2026にアクセス、https://medium.com/@david.j.sea/anthropic-just-passed-openai-in-revenue-here-is-why-it-matters-e3dd9bb04069
4. Anthropic $1 Trillion Secondary Valuation Overtakes OpenAI | LDS, 4月 26, 2026にアクセス、https://letsdatascience.com/blog/anthropic-1-trillion-secondary-market-overtakes-openai
5. Google to invest up to $40B in Anthropic as Claude outsells Gemini in the enterprise market it needs most - TNW, 4月 26, 2026にアクセス、https://thenextweb.com/news/google-40-billion-anthropic-investment-gemini
6. Anthropic raises $30 billion in Series G funding at $380 billion post-money valuation, 4月 26, 2026にアクセス、https://www.anthropic.com/news/anthropic-raises-30-billion-series-g-funding-380-b

illion-post-money-valuation
7. Anthropic revenue, valuation & funding - Sacra, 4月 26, 2026にアクセス、https://sacra.com/c/anthropic/
8. Claude AI Statistics and User Trends for 2026 | Second Talent, 4月 26, 2026にアクセス、https://www.secondtalent.com/resources/claude-ai-statistics/
9. Rewriting the IT Playbook: Empowering CIOs to Lead with Confidence in the AI Era, 4月 26, 2026にアクセス、https://techcommunity.microsoft.com/blog/microsoft365copilotblog/rewriting-the-it-playbook-empowering-cios-to-lead-with-confidence-in-the-ai-era/4411734
10. Microsoft in talks with Chevron, Engine No.1 over power deal- Bloomberg - Investing.com, 4月 26, 2026にアクセス、https://www.investing.com/news/stock-market-news/microsoft-in-talks-with-chevron-engine-no1-over-power-deal-bloomberg-4591912
11. Hallucination Stations - arXiv, 4月 26, 2026にアクセス、https://arxiv.org/pdf/2507.07505
12. Does the $O(N^2 \cdot d)$ computational complexity of Transformers pose a fundamental limit that makes hallucinations inevitable for complex tasks? - AI Stack Exchange, 4月 26, 2026にアクセス、https://ai.stackexchange.com/questions/50325/does-the-on2-cdot-d-computational-complexity-of-transformers-pose-a-fundam
13. Hallucination Stations: On Some Basic Limitations of Transformer-Based Language Models, 4月 26, 2026にアクセス、https://www.emergentmind.com/papers/2507.07505
14. AI (LLM) cannot operate beyond a certain level of complexity : r/ArtificialInteligence - Reddit, 4月 26, 2026にアクセス、https://www.reddit.com/r/ArtificialInteligence/comments/1qlrjw9/ai_llm_cannot_operate_beyond_a_certain_level_of/
15. Windows 10 ESU Cost in 2026: 3-Year Pricing Breakdown - Brytesoft, 4月 26, 2026にアクセス、https://brytesoft.com/blog/windows-10-esu-cost.html
16. Windows 10 Extended Security Updates (ESU): 2025 Pricing - IT GOAT, 4月 26, 2026にアクセス、https://www.itgoat.com/blog/windows-10-extended-security-updates-esu-2025-pricing/
17. Extended Security Updates (ESU) program for Windows 10 ..., 4月 26, 2026にアクセス、https://learn.microsoft.com/en-us/windows/whats-new/extended-security-updates
18. Prize for windows10 extension license - Microsoft Q&A, 4月 26, 2026にアクセス、https://learn.microsoft.com/en-us/answers/questions/5704668/prize-for-windows10-extension-license
19. How to Bypass SSE4.2 requirement on Windows 11 24H2? For learning and experimental purpose on old CPU such as Intel Core 2 Duo. - Reddit, 4月 26, 2026にアクセス、https://www.reddit.com/r/techsupport/comments/1g02xb9/how_to_bypass_sse42_requirement_on_windows_11/
20. Microsoft blocks some PCs from Windows 11 24H2 — CPU must support SSE4.2

or the OS will not boot | Tom's Hardware, 4月 26, 2026にアクセス、https://www.tomshardware.com/software/windows/microsoft-updates-windows-11-24h2-requirements-cpu-must-support-sse42-or-the-os-will-not-boot
21. Microsoft ending support for Windows 10 could send 240 mln PCs to landfills - report, 4月 26, 2026にアクセス、https://community.opentextcybersecurity.com/tech-talk-7/microsoft-ending-support-for-windows-10-could-send-240-mln-pcs-to-landfills-report-356456
22. Killing off Windows 10 could send 240million PCs to landfill - GB News, 4月 26, 2026にアクセス、https://www.gbnews.com/tech/windows-10-end-support-landfill
23. Microsoft's Windows 10 support shutdown may result in massive e-waste, warns Canalys Research | Domain-b.com, 4月 26, 2026にアクセス、https://www.domain-b.com/environment-1/microsoft-s-windows-10-support-shutdown-may-result-in-massive-e-waste-warns-canalys-research
24. Ending support for Microsoft's Windows 10 could lead to 240 million computers ending up in landfills - HackYourMom, 4月 26, 2026にアクセス、https://hackyourmom.com/en/novyny/prypynennya-pidtrymky-microsoft-windows-10-mozhe-pryzvesty-do-togo-shho-240-miljoniv-kompyuteriv-vykynutsya-na-zvalysha/
25. Planet Mozilla, 4月 26, 2026にアクセス、https://planet.mozilla.org/
26. Old habits die hard: Microsoft tries to limit our options, this time with AI | The Mozilla Blog, 4月 26, 2026にアクセス、https://blog.mozilla.org/en/mozilla/ai/microsoft-copilot-ai-user-choice/
27. Hallucination Stations: On Some Basic Limitations of Transformer-Based Language Models, 4月 26, 2026にアクセス、https://www.semanticscholar.org/paper/Hallucination-Stations%3A-On-Some-Basic-Limitations-Sikka-Sikka/fafc47ea54726e9dc0dccf948ba0292564880cc3
28. Memory Prices Surge Up to 90% From Q4 2025 - Counterpoint Research, 4月 26, 2026にアクセス、https://counterpointresearch.com/en/insights/Memory-Prices-Surge-Up-to-90-From-Q4-2025
29. Memory Prices Soar by 50% in Q4, Rally to Continue in 2026 - Counterpoint Research, 4月 26, 2026にアクセス、https://counterpointresearch.com/en/insights/memory-price-tracker-january-2026
30. Chevron eyes $7B Texas power plant for Microsoft data center campus, 4月 26, 2026にアクセス、https://energycapitalhtx.com/chevron-microsoft-power-plant
31. Chevron Could Build a $7 Billion Gas Plant to Power Microsoft's AI Ambitions. Time to Buy the Energy Giant? | The Motley Fool, 4月 26, 2026にアクセス、https://www.fool.com/investing/2026/04/02/chevron-could-build-a-7-billion-gas-plant-to-power/
32. Chevron, Microsoft Team Up for Giant Texas Gas Plant - RealClear Energy, 4月 26, 2026にアクセス、https://www.realclearenergy.org/2026/04/06/chevron_microsoft_team_up_for_giant_texas_gas_plant_1174109.html
33. Microsoft in Talks With Chevron, Engine No. 1 Over $7 Billion Texas Power Plant, 4月 26, 2026にアクセス、

https://energynow.com/2026/04/microsoft-in-talks-with-chevron-engine-no-1-over-7-billion-texas-power-plant/
34. Chevron statement regarding exclusivity agreement with Microsoft and Engine No. 1, 4月 26, 2026にアクセス、https://www.chevron.com/newsroom/2026/q1/chevron-statement-regarding-exclusivity-agreement-with-microsoft-and-engine-no-1
35. NYC Teachers' Retirement System Investment Meeting (Public), 4月 26, 2026にアクセス、https://comptroller.nyc.gov/wp-content/uploads/documents/04.16.2026-NYC-Teachers-Retirement-System-Investment-Meeting-Agenda-Package-Public.pdf
36. TRS Annual Climate Report FY2025 - Office of the New York City Comptroller Mark Levine, 4月 26, 2026にアクセス、https://comptroller.nyc.gov/reports/trs-annual-climate-report-fy2025/
37. Why Microsoft Is Investing $10 Billion in Japan? - TradingKey, 4月 26, 2026にアクセス、https://www.tradingkey.com/analysis/stocks/us-stocks/261750379-microsoft-japan-tradingkey
38. Microsoft's $10 Billion Japan Investment: AI Infrastructure, Sakura Internet, and Workforce Development Explained | Electronics Lifehack!, 4月 26, 2026にアクセス、https://electwork.net/en/news/microsoft-japan-10-billion-investment-ai-infrastructure-cybersecurity-workforce-2026/
39. Microsoft Bets $10 Billion to Boost Japan's AI, Cybersecurity - Dark Reading, 4月 26, 2026にアクセス、https://www.darkreading.com/cloud-security/microsoft-bets-10-billion-to-boost-japan-s-ai-cybersecurity
40. Microsoft: $10 Billion Investment To Expand AI Infrastructure In Japan - Pulse 2.0, 4月 26, 2026にアクセス、https://pulse2.com/microsoft-10-billion-investment-to-expand-ai-infrastructure-in-japan/
41. Microsoft Just Bet $10 Billion on Japan's AI Future — But the Country Still Can't Train Enough Engineers - Medium, 4月 26, 2026にアクセス、https://medium.com/@jinchannel6/microsoft-just-bet-10-billion-on-japans-ai-future-but-the-country-still-can-t-train-enough-d0ccee49dff6
42. Microsoft's $10B Japan AI Investment: SoftBank & Sakura - Tomorrow Access, 4月 26, 2026にアクセス、https://tomorrowaccess.com/en/news/18629
43. 2026 Iran war - Wikipedia, 4月 26, 2026にアクセス、https://en.wikipedia.org/wiki/2026_Iran_war
44. 2026 Iran war | Explained, United States, Israel, Strait of Hormuz, Map, & Conflict, 4月 26, 2026にアクセス、https://www.britannica.com/event/2026-Iran-war
45. Iran Update Evening Special Report, February 28, 2026 | ISW, 4月 26, 2026にアクセス、https://understandingwar.org/research/middle-east/iran-update-evening-special-report-february-28-2026/
46. Extended naval blockade is admission US military escalation poses even greater

risk, 4月 26, 2026にアクセス、
https://www.theguardian.com/world/2026/apr/23/extended-naval-blockade-is-an-admission-us-military-escalation-poses-even-greater-risk
47. live White House Says 'No Firm Deadline' For Ending War With Iran As US Naval Blockade Continues - Radio Free Europe, 4月 26, 2026にアクセス、
https://www.rferl.org/a/iran-war-us-hormuz-oil-blockade-gulf-israel/33640284/lbl0lbi444583.html
48. 2026 Strait of Hormuz crisis - Wikipedia, 4月 26, 2026にアクセス、
https://en.wikipedia.org/wiki/2026_Strait_of_Hormuz_crisis
49. Joint Force Enforces Maritime Blockade in Gulf of Oman, Globally, 4月 26, 2026にアクセス、
https://www.war.gov/News/News-Stories/Article/Article/4470128/joint-force-enforces-maritime-blockade-in-gulf-of-oman-globally/
50. Trump didn't extend Iran ceasefire. He launched an economic weapon, 4月 26, 2026にアクセス、
https://www.indiatoday.in/world/story/iran-war-trump-ceasefire-oil-fields-kharg-island-hormuz-blockade-2899805-2026-04-22